

VECTORS in R

Introduction to data science (DSC 105) Fall 2022

September 30, 2023

Contents



1 What will you learn?



After this lesson, you should be able to:

- Understand assignment in R
- Creating vectors, sequences and repetitions
- Sorting and measuring vector length
- Subsetting and extracting vector elements
- Vectorizing (rescaling)
- Classes and logical vectors

Most sections are accompanied by YouTube videos.

2 Everything is an object video



- R is a functional, object-oriented language
- There are usually many different ways to achieve the same result
- Different solutions differ in: performance, clarity, ease of use
- PERFORMANCE: important for large data sets only
- CLARITY: essential when communicating methods, code, results
- EASE OF USE: determines the fun you have when programming

[Watch YouTube Playlist Vectors Part I]

R is a functional, object-oriented language: everything's an object, and functions rule, as you already know. Because functions rule, there are usually many different ways to achieve the same result. They often differ in terms of performance, ease of use, and clarity. Performance becomes important when you work with truly large data sets, otherwise not so much. Ease of use to some extent determines the fun you do or don't have when using the language. Clarity is essential when communicating your methods (including your code) and your results to others. This is not a nice to have. Views of different packages (like the Tidyverse) differ massively regarding all of these.

3 Assigning objects (video)

- You can use `<-` or `=` for assignment of values to variables
- Use `<-` for objects, and `=` to assign function parameters

- Object names must start with a letter and avoid reserved words
- Challenge (15 min):

1. Create an object `foo` that stores the value $3^2 \times 4^{1/8}$ and display it.

```
foo <- 3**2 * 4**(1/8)
foo

[1] 10.70286
```

2. Overwrite `foo` with itself divided by 2.33. Print the result to the console.

```
foo <- foo / 2.33
foo

[1] 4.593504
```

3. Experiment with different formats and types - look all functions that you don't know up using `help`:

- (a) Create a new object `bar` with the value -8.2×10^{-13} using scientific ("e") notation and print it to the console.
- (b) Print `bar` again, but this time **without** scientific notation, using the `format` function. Save the result in `barf`
- (c) Check the data type of `barf` with the `class` function.
- (d) Convert `barf` back to a number with the `as.numeric` function.

```
bar <- -8.2e-13
bar
barf <- format(bar, scientific=FALSE)
barf
class(barf)
as.numeric(barf)

[1] -8.2e-13
[1] "-0.000000000000082"
[1] "character"
[1] -8.2e-13
```

4. Print the result of multiplying `foo` and `bar` to the console.

```
foo * bar

[1] -3.766673e-12
```

5. Check all variables created in the current session so far with `ls`, then delete them with `rm`, and check that they're all gone.

```
ls()
rm(list=ls())
ls()

[1] "bar"  "barf" "data" "foo"
character(0)
```

4 Why we need vectors (video)

- We need vectors to handle multiple items in R
- Vectors contain observations or measurements of a single variable
- Each vector can have only one data type - e.g. `numeric` ¹
- Give three examples of such collections for vectors:
 - numbers, e.g. the heights of students in this class
 - text, e.g. the names of students of this class
 - logical values, e.g. sex of students of this class
- Define sample vectors `s_heights`, `s_names` and `s_male` using the `c` function.
- Print the vectors.

Solution:

1. the heights of every student of this class, in cm.
2. the first names of every student of this class (strings).
3. observations, if a student is male or not male (male means `TRUE`).

¹Note: If a vector contains different data types, R coerces the vector elements to conform to one type, as we will see later. A data type that can hold any type of value is called a `list`.

Let's put some wood behind the arrow and define sample vectors for these:

```
s_heights <- c(180, 181, 158, 175, 179, 168)
s_names <- c("Vincent", "Natalija", "Adrian", "Andres", "Helena")
s_male <- c(TRUE, FALSE, TRUE, TRUE, FALSE)
s_heights
s_names
s_male

[1] 180 181 158 175 179 168
[1] "Vincent" "Natalija" "Adrian"   "Andres"   "Helena"
[1] TRUE FALSE TRUE TRUE FALSE
```

5 R object class



- The function `class` (check the help) returns the R object type
- Object type is not the same as data type or storage type
- Besides types there are also data structures (like `vector`)
- Let's check the class for our three vectors with student data

In order to check what R thinks about your observation or data type, you can use the function `class`. Look at the help page for details and enter the examples at the bottom of the help page.

Let's call `class` for our three sample vectors:

```
class(s_heights) # what type vector is this?
class(s_names)   # what type vector is this?
class(s_male)    # what type vector is this?

[1] "numeric"
[1] "character"
[1] "logical"
```

6 Creating vectors (video)

- The function to create a vector, or "combine values", is `c()`:

```
myvec <- c(1,3,1,42) # combine integers as vector
myvec                # prints 1 3 1 42
class(myvec)         # determine the data type - "numeric"

[1] 1 3 1 42
[1] "numeric"
```

- Vector elements can be calculations or previously stored items:

```
foo <- 32.1
myvec2 <- c(3, -3, 3.45, 1e+03, 64^0.5, 2+(3-1.1)/9.44, foo)
myvec2

[1] 3.000000 -3.000000 3.450000 1000.000000 8.000000 2.201271
[7] 32.100000
```

- Vector elements can even be vectors themselves:

```
myvec3 <- c(myvec, myvec2)
myvec3
class(myvec3)
```

```

[1] 1.000000 3.000000 1.000000 42.000000 3.000000 -3.000000
[7] 3.450000 1000.000000 8.000000 2.201271 32.100000
[1] "numeric"

```

- In the resulting output, the two vectors were put side by side. The new vector now has $11 = 4 + 7$ elements.
- What about missing values, `NA`, and non-numbers, `NaN`, and what about the special values `Inf` and `-Inf` - can you have these in vectors, too? Can you think about a way to test this?

```

specvec <- c(NA, NaN) # a vector with a NA and a NaN
class(specvec)
is.nan(specvec) # testing for NaN values
is.na(specvec) # testing for NA values
specvec1 <- c(specvec, Inf, -Inf) # a new vector with Inf, Inf
is.finite(specvec1) # testing for finiteness
is.infinite(specvec1) # testing for infiniteness
class(specvec1)

[1] "numeric"
[1] FALSE TRUE
[1] TRUE TRUE
[1] FALSE FALSE FALSE FALSE
[1] FALSE FALSE TRUE TRUE
[1] "numeric"

```

7 Base R - what's that again?

- Let's not forget that there are real data out there!
- R (i.e. "Base R") includes a number of built-in datasets

Base R Cheat Sheet

Getting Help

Accessing the help files

`?mean`
Get help of a particular function.

`help.search("weighted mean")`
Search the help files for a word or phrase.

`help(package = "dplyr")`
Find help for a package.

More about an object

`str(iris)`
Get a summary of an object's structure.

`class(iris)`
Find the class an object belongs to.

Using Libraries

`install.packages("dplyr")`
Download and install a package from CRAN.

`library(dplyr)`
Load the package into the session, making all its functions available to use.

`dplyr::select`
Use a particular function from a package.

`data(iris)`
Load a built-in dataset into the environment.

Working Directory

`getwd()`
Find the current working directory (where inputs are found and outputs are sent).

`setwd("C://file/path")`
Change the current working directory.

Use projects in RStudio to set the working directory to the folder you are working in.

Vectors

Creating Vectors

<code>c(2, 4, 6)</code>	<code>2 4 6</code>	Join elements into a vector
<code>2:6</code>	<code>2 3 4 5 6</code>	An integer sequence
<code>seq(2, 3, by=0.5)</code>	<code>2.0 2.5 3.0</code>	A numeric sequence
<code>rep(1:2, times=3)</code>	<code>1 2 1 2 1 2</code>	Repeat a vector
<code>rep(1:2, each=3)</code>	<code>1 1 1 2 2 2</code>	Repeat elements of a vector

Vector Functions

<code>sort(x)</code>	<code>rev(x)</code>
Return x sorted.	Return x reversed.
<code>table(x)</code>	<code>unique(x)</code>
See counts of values.	See unique values.

Selecting Vector Elements

By Position

<code>x[4]</code>	The fourth element.
<code>x[-4]</code>	All but the fourth.
<code>x[2:4]</code>	Elements two to four.
<code>x[-(2:4)]</code>	All elements except two to four.
<code>x[c(1, 5)]</code>	Elements one and five.

By Value

<code>x[x == 10]</code>	Elements which are equal to 10.
<code>x[x < 0]</code>	All elements less than zero.
<code>x[x %in% c(1, 2, 5)]</code>	Elements in the set 1, 2, 5.

Named Vectors

<code>x["apple"]</code>	Element with name 'apple'.
-------------------------	----------------------------

Programming

For Loop

```
for (variable in sequence){
  Do something
}
```

Example

```
for (i in 1:4){
  j <- i + 10
  print(j)
}
```

While Loop

```
while (condition){
  Do something
}
```

Example

```
while (i < 5){
  print(i)
  i <- i + 1
}
```

If Statements

```
if (condition){
  Do something
} else {
  Do something different
}
```

Example

```
if (i > 3){
  print("Yes")
} else {
  print("No")
}
```

Functions

```
function_name <- function(var){
  Do something
  return(new_variable)
}
```

Example

```
square <- function(x){
  squared <- x*x
  return(squared)
}
```

Reading and Writing Data

Input	Output	Description
<code>df <- read.table("file.txt")</code>	<code>write.table(df, "file.txt")</code>	Read and write a delimited text file.
<code>df <- read.csv("file.csv")</code>	<code>write.csv(df, "file.csv")</code>	Read and write a comma separated value file. This is a special case of read.table/write.table.
<code>load("file.Rdata")</code>	<code>save(df, file = "file.Rdata")</code>	Read and write an R data file, a file type special for R.

Conditions	a == b	Are equal	a > b	Greater than	a >= b	Greater than or equal to	is.na(a)	is missing
	a != b	NOT equal	a < b	Less than	a <= b	Less than or equal to	is.null(a)	is null

RStudio® is a trademark of RStudio, Inc. • © 2018 Robert I. Kabacoff • r.i.kabacoff@gmail.com

Learn more at [web page](#) or [vignette](#) • package version • Updated 3/18

- Do you remember how to find these pre-loaded datasets? One of these is `Nile`. Do you remember how to get information on such a dataset (or on any R object)?

9

8 Example: down the Nile (video)



- What data structure is `Nile`?
- `Nile` contains a so-called "time series", a sequence of numbers that correspond to measurements of the annual flow (in billion - 10 - cubic meters) of the river Nile at Aswan, measured between 1871-1970. You can use `class` to confirm it:

```
class(Nile) # what type of dataset is this?  
  
[1] "ts"
```

The output is `"ts"` or time series. You may remember that we previously looked at large datasets. `mtcars` for example was a "data frame" (we'll learn more about them later).

- How can we print this dataset, or parts of it, on the screen? **Solution:** there are different ways to look inside `Nile`:

```
str(Nile) # show dataset structure  
head(Nile) # show first few elements  
Nile      # this prints the whole dataset
```

```

Time-Series [1:100] from 1871 to 1970: 1120 1160 963 1210 1160 1160 813 1230 1370
[1] 1120 1160 963 1210 1160 1160
Time Series:
Start = 1871
End = 1970
Frequency = 1
[1] 1120 1160 963 1210 1160 1160 813 1230 1370 1140 995 935 1110 994 1020
[17] 1180 799 958 1140 1100 1210 1150 1250 1260 1220 1030 1100 774 840 874
[33] 940 833 701 916 692 1020 1050 969 831 726 456 824 702 1120 1100
[49] 764 821 768 845 864 862 698 845 744 796 1040 759 781 865 845
[65] 984 897 822 1010 771 676 649 846 812 742 801 1040 860 874 848
[81] 744 749 838 1050 918 986 797 923 975 815 1020 906 901 1170 912
[97] 919 718 714 740

```

Because we don't know yet how to look at sub-vectors or individual vector elements, we cannot directly check what type the elements of `Nile` have, but the output seems to suggest that the Nile flow is measured in integer numbers.

You can also see from the print output of `Nile` how row labels work: there are 15 numbers per row, and the second row starts with the 16th number, indicated by `[16]`.

9 Plotting the nile (video)

- Plotting is often a good entry into exploring data
- `Nile` is a numeric vector of a single, continuous variable over time
- To visualize such data, *histograms* or *line plots* are useful
- What you're really after is a picture of a value *distribution*
- Why are *histograms* called "histograms"?
- How can you find out more about plotting a histogram in R?

HELP: You know of course what to do at this point: call for help using `?hist`. Skip to the **Examples** section at the end, where you find the command `hist(islands)`. This creates a histogram of another dataset, `islands`. With the help of `?islands`, you find out quickly that this is a "named vector of 48 elements". Never mind what this


means, but you can enter the command, which will generate a plot. This is a histogram: it plots frequency of the data and distributes them into bins². Let's get back to the river Nile.

Like most R functions, `hist` has many options. If you execute `hist(Nile)`, you get the same type of graph as in the example except that we know what the data are (annual Nile flow measurements in 10 m³, or 100,000,000 (100 million) of cubic metres).

10 Plotting the histogram

- Let's plot the histogram of Nile


```
hist(Nile)
```



²The Wikipedia entry for "histogram" is not bad as a start, lots of examples and you'll soon find out how to make these yourself! The origin of the name "histogram" is not clear - it was probably invented by Pearson, who introduced this type of graph, and is short for "HISTorical diaGRAM".

- ☐ Can you interpret the plot given what you know about the data?
- ☐ Add the argument `breaks=20` to the `hist` function call. Change the file name in the code block header if you want to create a new PNG file

```
hist(Nile, breaks=20)
```



The `hist` function creates 10 bins by default and distributes the data accordingly. You can alter this number of bins by changing the argument `breaks`, e.g. `hist(Nile, breaks=20)` (try it!).


We'll get back to the `Nile` once we know more about vectors! In the next four sections, we're going to look at useful functions.

11 Plotting the line plot

Since `Nile` is a time series, every data point has a time label

You can easily plot the evolution of the data over time with `plot`
A line plot is useful to visualize two continuous numeric variables
This leads to a so-called *line plot*

```
plot(Nile)
```




☐ Can you interpret the plot given what you know about the data?

Is `hist` a generic function?

- Try to make a histogram of `mtcars`.
- Check out the structure of `mtcars` to see the column vector data types.
- Make a histogram of `mtcars$mpg`:

```
hist(mtcars$mpg)
```



Will `hist` work with `factor` vectors? Check out `ToothGrowth`, another built-in dataset:

```
str(ToothGrowth)
```

```
'data.frame': 60 obs. of 3 variables:
 $ len : num  4.2 11.5 7.3 5.8 6.4 10 11.2 11.2 5.2 7 ...
 $ supp: Factor w/ 2 levels "OJ","VC": 2 2 2 2 2 2 2 2 2 2 ...
 $ dose: num  0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 ...
```

Try to create a histogram of the `supp` `factor` vector:

```
tgs <- ToothGrowth$supp
```

```
# hist(tgs) # Error in hist.default(tgs) : 'x
```

Can we turn the `factor` values into numbers?

```

tgs <- as.numeric(tgs)
tgs

[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1 1 1
[41] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1


```

Now we can plot the values as a histogram:

```

hist(tgs,
      breaks = 2,
      col = "purple",
      border = "white",
      xlab = "Supply type (orange juice/vitamin C)")

```



For these and other settings, check out `help(hist)`, and `example(hist)` for the examples listed at the end of the documentation (see next section).

12 Asking for help



- When you see a new function or dataset, look it up
- Use fuzzy help search (??) or regular help (?, help)
- Scroll down to check out (and run) the **examples**
- Get an overview of the available options

In the following, I won't waste more space with the obvious: whenever I mention a new function or dataset, or keyword, look the corresponding help up immediately. More often than not, you will take something away from it - at the very minimum an example. Over time, you'll understand things even though you don't know how you possibly could: this is because you've begun to develop a habit by using a system of learning - looking up the help content - and the more you look at help pages, the more you recognize known concepts.

13 Creating sequences and repetitions (video)



14 The colon : operator

`1:n` creates a sequence of numbers separated by intervals of 1

`3:21`

```
[1] 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21
```

Check what type of R object `3:21` is by applying the functions:

- `class` (R object class),
- `mode` (R object storage mode)
- `is.vector` (R vector check)

```
class(3:27)
mode(3:27)
is.vector(3:27)
```

```
[1] "integer"
[1] "numeric"
[1] TRUE
```

Sequences created this way can also be stored.

```
foo <- 5.3
bar <- foo:10
bar
```

```
[1] 5.3 6.3 7.3 8.3 9.3
```

What happens if the first argument of `:` is smaller than the second?

```
x <- 10:foo
x
```

```
[1] 10 9 8 7 6
```

You can perform computations to specify the range.

```
baz <- foo:(-47+1.5)
baz
```


```
[1] 5.3 4.3 3.3 2.3 1.3 0.3 -0.7 -1.7 -2.7 -3.7 -4.7 -5.7 -6.7
[14] -7.7 -8.7 -9.7 -10.7 -11.7 -12.7 -13.7 -14.7 -15.7 -16.7 -17.7 -18.7 -19.7
[27] -20.7 -21.7 -22.7 -23.7 -24.7 -25.7 -26.7 -27.7 -28.7 -29.7 -30.7 -31.7 -32.7
[40] -33.7 -34.7 -35.7 -36.7 -37.7 -38.7 -39.7 -40.7 -41.7 -42.7 -43.7 -44.7
```

Try to understand what happened here by checking the numbers: the first value of the sequence is `foo = 5.3`. The last value is a negative value, `-47+1.5 = -45.5`. In order to generate the sequence, R counts down in steps of 1 from the first to the last value. It stops at `-44.7`, because the next value, `-45.7` would be outside of the interval `[5.3, -45.5]`.

14.1 Challenge: weird histogram

The following code creates a weird histogram:

```
hist(1:10, breaks=10)
```



What can you do to see the proper distribution with breaks at every number (all bars should have height 1)?

15 Sequences

The function `seq` allows modifying the step-width with `by`:

```
seq(from = 3, to = 27, by = 3)
```

```
[1]  3  6  9 12 15 18 21 24 27
```

`seq` always starts at `from` but not always end on `to`:

```
seq(from=1, to=10, by=2) # range even, stepsize even  
seq(from=1, to=11, by=2) # range odd, stepsize even
```

```
[1] 1 3 5 7 9
[1] 1 3 5 7 9 11
```

To end exactly on the last value, use `length.out`:

```
seq(from=1, to=10, length.out=10) # either by or length.out
seq(from = 3, to = 27, length.out = 40)
```

```
[1] 1 2 3 4 5 6 7 8 9 10
[1] 3.000000 3.615385 4.230769 4.846154 5.461538 6.076923 6.692308 7.307692
[9] 7.923077 8.538462 9.153846 9.769231 10.384615 11.000000 11.615385 12.230769
[17] 12.846154 13.461538 14.076923 14.692308 15.307692 15.923077 16.538462 17.153846
[25] 17.769231 18.384615 19.000000 19.615385 20.230769 20.846154 21.461538 22.076923
[33] 22.692308 23.307692 23.923077 24.538462 25.153846 25.769231 26.384615 27.000000
```

What is the step-width in the last case? Compute it and use it to create a sequence of 40 numbers from 3 to 27 exactly, with `seq`.

```
s <- seq(from = 3, to = 27, length.out = 40)
s[2]-s[1] # step-width
seq(from = 3, to = 27, by = s[2]-s[1])
```

```
[1] 0.6153846
[1] 3.000000 3.615385 4.230769 4.846154 5.461538 6.076923 6.692308 7.307692
[9] 7.923077 8.538462 9.153846 9.769231 10.384615 11.000000 11.615385 12.230769
[17] 12.846154 13.461538 14.076923 14.692308 15.307692 15.923077 16.538462 17.153846
[25] 17.769231 18.384615 19.000000 19.615385 20.230769 20.846154 21.461538 22.076923
[33] 22.692308 23.307692 23.923077 24.538462 25.153846 25.769231 26.384615 27.000000
```

`length.out` can only be positive (there is no 'negative length').

Create a decreasing sequence of length 5 from 5 to -5. Use `length.out` first, then use `by` to achieve the same result.


```
myseq <- seq(from=5, to=-5, length.out=5)
myseq
myseq1 <- seq(from=5, to=-5, by = -2.5)
myseq1
```

```
[1] 5.0 2.5 0.0 -2.5 -5.0
[1] 5.0 2.5 0.0 -2.5 -5.0
```

16 Odd histogram with :

When creating a histogram of the vector `1:10` with binwidth 10, the resulting graph looks wrong:

```
hist(1:10, breaks=10)
```

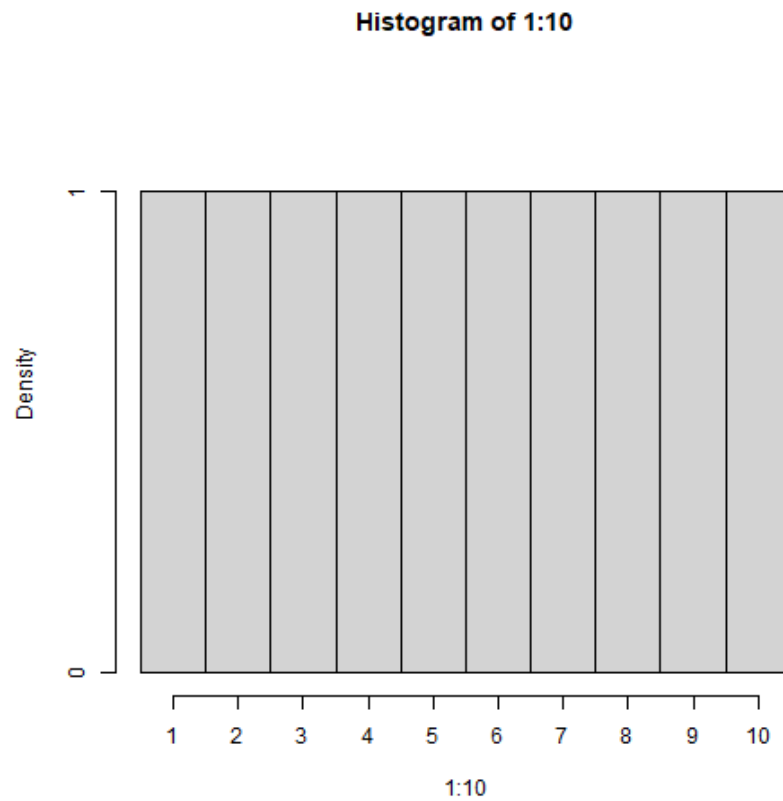


The code below fixes all problems:

1. `freq = FALSE` normalizes the diagram so that it shows the density (number of elements over total number of elements) instead of count.
2. `breaks = seq(0.5,10.5,by=1)` puts each integer in its own bin.
3. `axes = FALSE` removes the axes so that we can redraw them.
4. `ylim = c(0,0.12)` resets the y-axis values (for the density $\in [0,0.1]$).

5. The `axis` commands redraw the axis, for the y-axis including labels.

```
hist(1:10,  
     freq = FALSE,  
     breaks = seq(0.5, 10.5, by=1),  
     axes = FALSE,  
     ylim = c(0, 0.12))  
axis(1, at = 1:10)  
axis(2, at = c(0,0.1), labels = c('0','1'))
```



17 NEXT Repetition

- Use `rep` to repeat a value, e.g. the number 1, four times:

```
rep( x = 1, times = 4)
```

```
[1] 1 1 1 1
```

- You can repeat any R object, e.g. the vector `c(3, 62, 8, 3)`, or the scalar `foo`, or an arbitrary arithmetic expression:

```
rep(x=c(3,62,8,3),times=3)
rep(x=foo, times=2)
rep(x=2*2, times=(foo*2)) # times must be a positive integer
```

```
[1] 3 62 8 3 3 62 8 3 3 62 8 3
[1] 5.3 5.3
[1] 4 4 4 4 4 4 4 4 4 4
```

- The `each` argument of `rep(x)` says how many times each element of `x` is repeated:

```
rep(x=c(3,62), times=3) # repeat vector three times
rep(x=c(3,62), each=2)  # repeat each vector element twice
rep(x=c(3,62), times=3, each=2) # repeat each vector element twice,
                                # and repeat the result three times
```

```
[1] 3 62 3 62 3 62
[1] 3 3 62 62
[1] 3 3 62 62 3 3 62 62 3 3 62 62
```

☐ The default of `times` and `each` is 1. What is `rep(c(3,62))` ?

- As with `seq`, you can include the result of `rep` in a vector of the same data type (e.g. `numeric`):

```
foo <- 4 # store 4 in foo

## create vector with rep and seq
c(3, 8.3, rep(x=32,times=foo), seq(from=-2,to=1,length.out=foo+1))

[1] 3.00 8.30 32.00 32.00 32.00 32.00 -2.00 -1.25 -0.50 0.25 1.00
```


18 Repetition with characters

- `rep` also works for characters and character vectors:

```
rep(x="data science", times=2)
rep(x=c("data","science"), times=2)
rep(x=c("data","science"), times=2, each=2)
```

```
[1] "data science" "data science"
[1] "data"      "science" "data"      "science"
[1] "data"      "data"      "science" "science" "data"      "data"      "science" "science"
```

- What happens if you try to mix characters, numbers, Booleans? Repeat an expression that has all three data types in it.

```
rep(x=c("data", 1, TRUE), times=2)
```

```
[1] "data" "1"      "TRUE" "data" "1"      "TRUE"
```

When you call a function with an argument of the wrong type, or, as in the case of `c`, you try to create a vector of different data types, R responds with "coercion" to make it happen.

19 Sorting and measuring lengths (video)



20 Sorting vector elements

- *Sorting* is important because we don't care about memory locations
- `sort(x)` arranges the elements of `x` according to size
- The default order is ascending, or `decreasing = FALSE`

```
sort(x = c(2.5, -1, -10, 3.44)) # sort ascending
sort(x = c(2.5, -1, -10, 3.44), decreasing = FALSE) # sort ascending
sort(x = c(2.5, -1, -10, 3.44), decreasing = TRUE) # sort descending
```

```
[1] -10.00 -1.00  2.50  3.44
[1] -10.00 -1.00  2.50  3.44
[1]  3.44  2.50 -1.00 -10.00
```

- Special values are removed, put last or first with `na.last`. This works for all special values - NA, NaN and Inf.

```
sort(x = c(2.5, -1, -10, 3.44, NA), na.last=TRUE) # put NA last
sort(x = c(2.5, -1, -10, 3.44, NaN), na.last=TRUE) # put NaN last
sort(x = c(2.5, -1, -10, 3.44, Inf), na.last=TRUE) # put Inf last
sort(x = c(2.5, -1, -10, 3.44, NA), na.last=FALSE) # put NA first
sort(x = c(2.5, -1, -10, 3.44, NA), na.last=NA) # remove NA
```

```
[1] -10.00 -1.00  2.50  3.44    NA
[1] -10.00 -1.00  2.50  3.44   NaN
[1] -10.00 -1.00  2.50  3.44    Inf
[1]      NA -10.00 -1.00  2.50  3.44
[1] -10.00 -1.00  2.50  3.44
```

- Remember that NA is a logical object. How can you check that?

```
class(NA)
```

```
[1] "logical"
```

21 Length of vectors

- The `length` function gets or sets the length of vectors³:

```
length(x = c(3,2,8,1,10)) # vector of 5 elements
length(x = 5:13)          # vector of 9 elements
length(x = c(3,2,2^3,5*3)) # vector of 4 elements
length(1000)              # scalar/vector of 1 element

[1] 5
[1] 9
[1] 4
[1] 1
```

- If you have functions inside the object definition, `length` gives you the number of entries *after* the inner functions have been executed:

```
foo <- 4
bar <- c(3,8.3,rep(x=32,times=foo),seq(from=-2,to=1,length.out=foo+1))
bar
length(bar)

[1] 3.00 8.30 32.00 32.00 32.00 32.00 -2.00 -1.25 -0.50 0.25 1.00
[1] 11
```

- R's display options are stored in `options()`, which is a `list`. Lists have a length like options. How many options does `options()` have?

```
class(options()) # class of options()
length(options()) # length of options() : number of options
class(options)   # class of options as a function
class(options())$digits # class of one options() element
length(options())$digits # length of one options() element

[1] "list"
[1] 66
[1] "function"
[1] "integer"
[1] 1
```

³Both `length` and `sort`, as you can read in the respective help pages, work both for vectors and for factors. These are necessary whenever we deal with qualities or categories (like "male" or "female") rather than quantities. You'll learn about them soon!

22 Practice: creating vectors



- Practice what you've learnt by solving problems independently.
- You find the practice file in here in DataCamp workspace.

23 Naming vectors

- Naming vector elements makes code more readable.

```
c(apple = 1, banana = 2, "kiwi fruit" = 3, 4)
```

```
apple      banana kiwi fruit
      1         2         3         4
```

- Or you can name elements explicitly using the function `names`⁴

```
x <- 1:4
names(x) <- c("apple", "bananas", "kiwi fruit", "")
x
names(x)
```

⁴You should look up the examples in `help(names)`: the data set `islands` is a named vector suited to play around with vector naming.

```

      apple      bananas kiwi fruit
      1          2          3          4
[1] "apple"      "bananas"    "kiwi fruit" ""

```

- Looking under the hood of `names`:

```

foo <- 1:4 # vector 1,2,3,4
names(foo) # vector is not named (NULL)
names(foo) <- letters[1:4] # assign letter names
names(foo)
foo # default display includes names
str(foo) # structures reveals names as attributes
attributes(foo) # attributes is a list of 1 element, $names
str(attributes(foo))

```

```

NULL
[1] "a" "b" "c" "d"
a b c d
1 2 3 4
Named int [1:4] 1 2 3 4
- attr(*, "names")= chr [1:4] "a" "b" "c" "d"
$names
[1] "a" "b" "c" "d"
List of 1
 $ names: chr [1:4] "a" "b" "c" "d"

```

24 Length of names vs. vector

- What if your names are too short (or too long) for your vector?
- Define a vector `week` whose elements are the names of weekdays

```

week <- c("Mon", "Tue", "Wed", "Thu", "Fri", "Sat", "Sun")
week

```

```

[1] "Mon" "Tue" "Wed" "Thu" "Fri" "Sat" "Sun"

```

- Define a vector `foo` that contains seven numbers, and name its elements according to the `week`.

```
foo <- 1:7
names(foo) <- week
foo
```

```
Mon Tue Wed Thu Fri Sat Sun
  1   2   3   4   5   6   7
```

- Copy `week` to `workweek`, and remove "Saturday" and "Sunday" from the `workweek`.

```
workweek <- week[1:5] # or week[-(6:7)]
workweek
```

```
[1] "Mon" "Tue" "Wed" "Thu" "Fri"
```

- Copy `foo` to `bar`, and overwrite `names(bar)` with `workweek`.

```
bar <- foo
bar
names(bar) <- workweek
names(bar[6:7]) # names of the last two elements missing - NA
names(bar) <- NULL # remove names altogether
bar
```

```
Mon Tue Wed Thu Fri Sat Sun
  1   2   3   4   5   6   7
[1] NA NA
[1] 1 2 3 4 5 6 7
```

25 Indexing vectors

- Passing a vector of positive numbers returns the slice of the vector containing the elements at those locations.

```
x <- (1:5)^2 # example vector
x
x[1] # extract the first element only
x[c(1,3,5)] # extract elements with indices 1,3,5
```

```
[1] 1 4 9 16 25
[1] 1
[1] 1 9 25
```

- Passing a vector of negative numbers returns the slice of the vector containing the elements everywhere except at those locations.

```
x[c(-2,-4)]
```

```
[1] 1 9 25
```

- Passing a logical vector returns the slice of the vector containing the elements where the index is TRUE.

```
x[c(TRUE, FALSE, TRUE, FALSE, TRUE)]
```

```
[1] 1 9 25
```

- For named vectors, passing a character vector of names returns the slice of the vector containing the elements with those names.

```
names(x) <- c("one", "four", "nine", "sixteen", "twenty five")
x[c("one", "nine", "twenty five")]
```

```
one          nine twenty five
 1              9          25
```

26 Coercion

- All vector elements have to be of the same `class` or type
- When you try to mix them, R will create vectors with "coercion":

```
foo <- c("a",NA,1)
foo
class(foo) # foo becomes a character vector
```

```
[1] "a" NA "1"
[1] "character"
```

- Missing values `NA` are not coerced to `character` (e.g. `"NA"`) because this would mean altering their main property, to be missing.
- Still, the whole vector is a `character` vector object:

```
mode(foo) # R storage mode
class(foo) # R object class
```

```
[1] "character"
[1] "character"
```

- You can also explicitly convert elements using the functions `as.character`, `as.logical` or `as.numeric`.

```
as.character(c(1,2,TRUE)) # convert vector to character values
as.numeric(c("a",2,TRUE)) # R turns characters and Booleans into NA
as.logical(c("a",0,TRUE)) # R turns characters and numerics into NA
```

```
[1] "1" "2" "1"
[1] NA  2 NA
Warning message:
NAs introduced by coercion
[1]  NA  NA TRUE
```

- Conversion with `as.logical` has a surprise: any non-zero number is turned into `TRUE` if the vector is `numeric`.

```
as.logical(c(1,0,-1, 0.333, -Inf, NaN ))
```

```
[1]  TRUE FALSE  TRUE  TRUE  TRUE    NA
```

- The lesson: don't mix data types in vectors if you can avoid it!

27 Summary with examples

- R is a functional language in which everything's an object.
- R functions differ in: performance (speed), ease-of-use and clarity.
- To assign values to objects, use the `<-` operator.

- To assign values to arguments in functions, use the `=` operator.
- The elements of a numeric, character or logical vector are numbers, letters or truth values.
- A vector can have arithmetic calculations or vectors as elements.
- A histogram distributes data by frequency across evenly spaced bins.
- Sequences of numbers can be created using the colon operator, or the functions `seq` or `rep`.
- Vectors can be sorted with `sort` in either direction.
- Vector length can be measured as the number of vector elements with `length`.
- Index vectors can be used to select sub-vectors.
- Negative index values delete the corresponding vector elements

R CODE EXAMPLES:

<code>x <- 5</code>	assign 5 to object <code>x</code>
<code>x <- x+1</code>	overwrite <code>x</code> (new value)
<code>c(1,2,3,4)</code>	define (numerical) vector
<code>class(bar)</code>	check type of object <code>bar</code>
<code>hist(x,breaks=foo)</code>	histogram of dataset <code>x</code> with <code>foo</code> bins
<code>m:n</code>	sequence <code>m</code> to <code>n</code> at intervals = 1
<code>seq(from=foo,to=bar,by=baz)</code>	sequence from <code>foo</code> to <code>bar</code> intervals = <code>baz</code>
<code>seq(from=foo,to=bar,length.out=fuz)</code>	seq. <code>foo</code> to <code>bar</code> , <code>fuz</code> equal intervals
<code>rep(x=foo,times=bar,each=baz)</code>	repeat <code>foo</code> times <code>bar</code> , and repeat each element of <code>foo</code> times <code>baz</code>
<code>vector("numeric",foo), numeric(foo)</code>	empty numeric vector of length <code>foo</code>
<code>vector("character",foo), character(foo)</code>	empty character vector of length <code>foo</code>
<code>vector("logical",foo), logical(foo)</code>	empty logical vector of length <code>foo</code>
<code>sort(x=foo, decreasing=FALSE)</code>	sort vector <code>foo</code> from smallest to largest
<code>sort(x=foo, decreasing=TRUE)</code>	sort vector <code>foo</code> from largest to smallest
<code>length(x=foo)</code>	print length of vector <code>foo</code>
<code>[n], [n:m], [-n]</code>	indices <code>n</code> , <code>n</code> to <code>m</code> , deleting element <code>n</code>
<code>prod(foo), sum(foo)</code>	multiply / sum up all elements of vector <code>foo</code>
<code>names(x)</code>	return names of vector <code>x</code> (or NULL)
<code>as.character, as.numeric, as.logical</code>	coerce arguments to the resp. class

28 Concept summary



- In R mathematical expressions are evaluated according to the *PEMDAS* rule.
- The natural logarithm $\ln(x)$ is the inverse of the exponential function e^x .
- In the scientific or e-notation, numbers are expressed as positive or negative multiples of 10.
- Each positive or negative multiple shifts the digital point to the right or left, respectively.
- Infinity `Inf`, not-a-number `NaN`, and not available numbers `NA` are *special values* in R.

29 Code summary

CODE	DESCRIPTION
<code>log(x=,b=)</code>	logarithm of <code>x</code> , base <code>b</code>
<code>exp(x)</code>	e^x , exp[onential] of x
<code>is.finite(x)</code>	tests for finiteness of <code>x</code>
<code>is.infinite(x)</code>	tests for infiniteness of <code>x</code>
<code>is.nan(x)</code>	checks if <code>x</code> is not-a-number
<code>is.na(x)</code>	checks if <code>x</code> is not available
<code>all.equal(x,y)</code>	tests near equality
<code>identical(x,y)</code>	tests exact equality
<code>1e2, 1e-2</code>	$10^2 = 100$, $10^{-2} = \frac{1}{100}$

30 References

- Richard Cotton (2013). Learning R. O'Reilly Media.
- Tilman M. Davies (2016). The Book of R. (No Starch Press).
- Rafael A. Irizarry (2020). Introduction to Data Science (also: CRC Press, 2019).
- Norman Matloff (2020). faster: Fast Lane to Learning R!.