

# Basic plotting in R - histograms, boxplots, scatterplots

Introduction to data science (DSC 105) Fall 2022

## Table of Contents

- [1. Histograms](#)
- [2. Histogram analysis](#)
- [3. Histograms in ggplot2](#)
- [4. Practice creating and customizing a histogram](#)
- [5. Boxplots - global summary stats](#)
- [6. Side-by-side boxplots](#)
- [7. Scatterplots](#)
- [8. Scatterplots of more than two variables](#)
- [9. References](#)

## 1 Histograms

- [ ]

When do you think of using **barplots**?

Barplots are sensible for counting **observations of categories**



- [ ]

When do you think of using a **histogram**?

When you have observed a single **numeric-continuous variable**

- *Example:* You've already seen the histogram for a simple data structure, the **time series** `ts` (here next to the **line plot**):

```
par(mfrow=c(2,1)) # create 2 x 1 plot array
plot(Nile)
hist(Nile)
```



- [ ]

What does the *height* of a bar represent exactly?

The height of each bar (on the y-axis) represents the number of years in which the volume of water flowing through the Nile was within its interval of 100 mio cubic metres width (on the x-axis).

- [ ]

How can you find out what the *binwidth* of this histogram is?


```
str(hist(Nile))
```

```
List of 6
 $ breaks  : int [1:11] 400 500 600 700 800 900 1000 1100 1200 1300 ...
 $ counts  : int [1:10] 1 0 5 20 25 19 12 11 6 1
 $ density : num [1:10] 0.0001 0 0.0005 0.002 0.0025 0.0019 0.0012 0.0011 0.0006 0.0001
 $ mids    : num [1:10] 450 550 650 750 850 950 1050 1150 1250 1350
 $ xname    : chr "Nile"
 $ equidist: logi TRUE
 - attr(*, "class")= chr "histogram"
```

- [ ]

Look at the help for `hist`, then change the binwidth to 30.

```
hist(Nile, breaks=30)
```



- [ ]

How can you **print** the value for the binwidth (breaks)?

```
h <- hist(Nile)
h$breaks[2]-h$breaks[1]
```

```
[1] 100
```

- [ ]

What happens when you choose `breaks=0`?

```
hist(Nile, breaks=0)
```

```
Error in hist.default(Nile, breaks = 0) : invalid number of 'breaks'
```

- [ ]

What happens when you set `breaks=1000001` ( $1e+6 + 1$ )?

```
hist(Nile, breaks=1000001)
```

```
Warning message:
In hist.default(Nile, breaks = 1000001) :
  'breaks = 1e+06' is too large and set to 1e6
```

- [ ]

Create a script file `Nile.R`, put the previous command into it, and run it as a batch process in the shell (using `M-x eshell`):

```
$ R CMD BATCH Nile.R
```

Open the output file `Nile.rout` to see the result:

```
> hist(Nile, breaks=1000001)
Warning message:
In hist.default(Nile, breaks = 1000001) :
  'breaks = 1e+06' is too large and set to 1e6
>
> proc.time()
   user  system elapsed 
  1.09    0.09    1.20
```

- `proc.time` and `system.time` measure the performance of R. Example: measure rolling 1 die a million times:

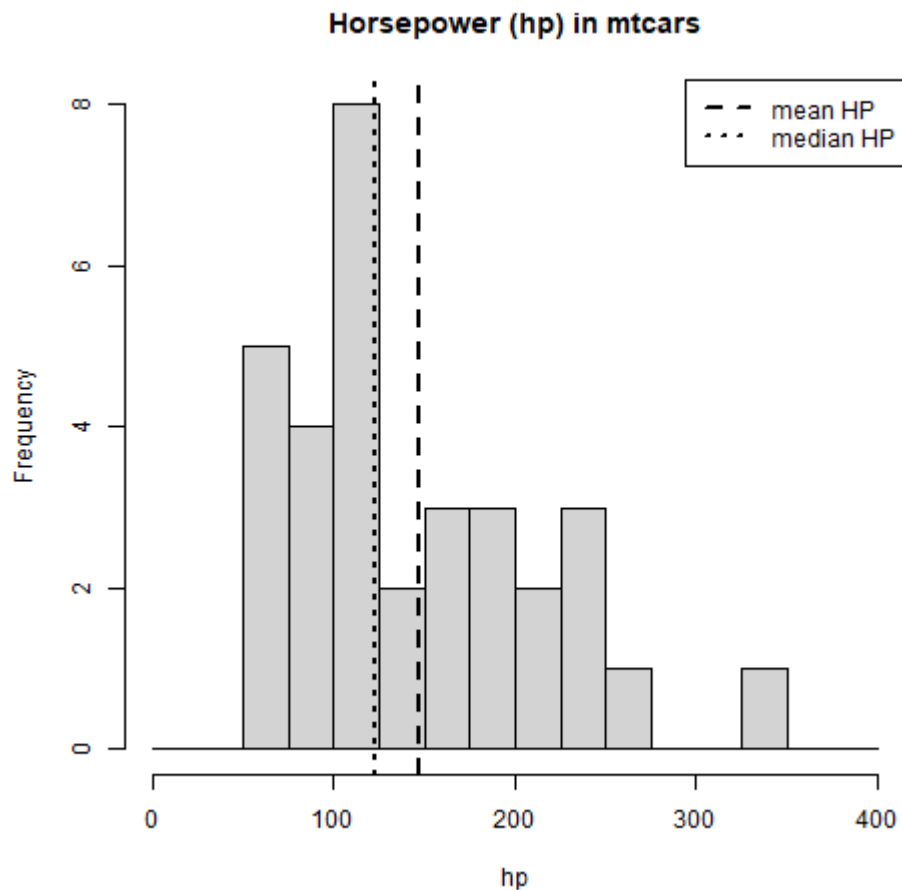
```
system.time(
  for (i in 1:1e6)
    sample(1:6,1))
```

```
user  system elapsed 
4.29    0.51    4.83
```

## 2 Histogram analysis

- You can manually set the histogram breaks by supplying a vector
- Example: horsepower `hp` in the `mtcars` dataset with breaks from 0 to 400, 25 units apart from each other:

```
hp <- mtcars$hp
hist(x=hp,
     breaks=seq(from=0, to=400, by=25),
     main="Horsepower (hp) in mtcars") # histogram
abline(v=c(mean(hp),median(hp)),
       lty=c(2,3),lwd=2) # mean, median
legend("topright",
       legend=c("mean HP","median HP"),
       lty=c(2,3),lwd=2) # Legend
```



- Reducing the bin width allows seeing more detail but also risks highlighting irrelevant features (like the single outlying car).

```
max(hp) # outlier in mtcars$hp
```


```
[1] 335
```

- Too small a binwidth leads to too much detail
- Too large a binwidth leads to loss of detail

### 3 Histograms in ggplot2


- Histogram of `mtcars$hp` in `ggplot2`:

```
library(ggplot2)
ggplot(data=mtcars,
       aes(x=hp)) +
  geom_histogram(binwidth=25, color="white", fill="red")
```



- A histogram of Nile is not so easy because it's a time series:

```
library(ggplot2)
ggplot(data=data.frame(Nile),
       aes(x=Nile)) +
  geom_histogram(color="black", fill="green")
```



## 4 Practice creating and customizing a histogram



- Download and open in GNU Emacs: [tinyurl.com/nhkykkxr](https://tinyurl.com/nhkykkxr)
- Work with a friend or with your neighbor in class
- Complete the problems in the file (we'll discuss at the end)
- Upload the result (for each participants) [to Canvas \(Practice 10\)](#)

## 5 Boxplots - global summary stats

- Open the practice file in Emacs to code along: [tinyurl.com/2e6dy9yb](https://tinyurl.com/2e6dy9yb)
- Box-and-whisker plots, or boxplots represent the five-number summary:
  1. Minimum
  2. 1st quartile (25% of the values are below it)
  3. Median (50% of the values are below/above it)
  4. Mean (Arithmetic average)
  5. 3rd quartile (75% of the values are below it)
  6. Maximum
- For example for `x <- c(1,2,3,4,5,6,7,8,9,10)`:

```
x <- c(1:10)
summary(x)
```

```
Error: unexpected symbol in:
"x <- c(1:10
summary"
```

- For the built-in quakes data frame of 1,000 seismic events near the island of Fiji (depth, magnitude, number of observing stations):

```
summary(quakes[,c("depth", "mag", "stations")])
```

depth	mag	stations
Min. : 40.0	Min. : 4.00	Min. : 10.00
1st Qu.: 99.0	1st Qu.: 4.30	1st Qu.: 18.00
Median : 247.0	Median : 4.60	Median : 27.00
Mean : 311.4	Mean : 4.62	Mean : 33.42
3rd Qu.: 543.0	3rd Qu.: 4.90	3rd Qu.: 42.00
Max. : 680.0	Max. : 6.40	Max. : 132.00

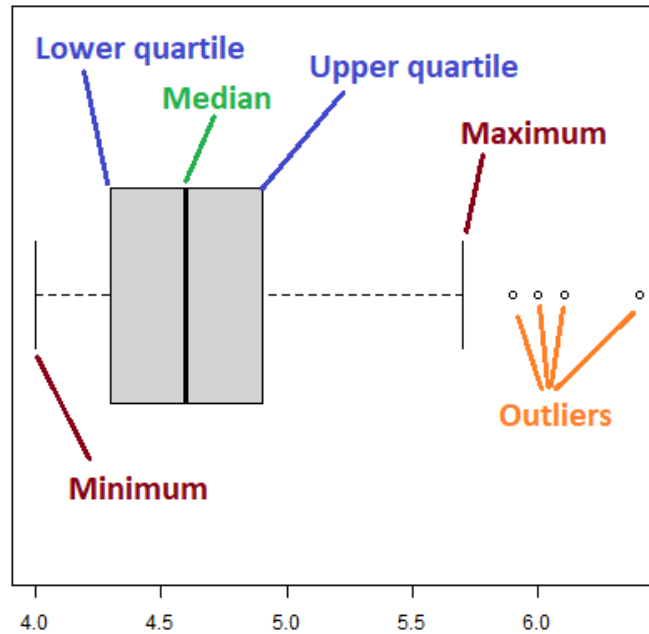
- Each record/row represents one recorded earthquake: where it was located, at which depth the epicenter was, its magnitude, and the number of observing stations:

```
head(quakes)
```

	lat	long	depth	mag	stations
1	-20.42	181.62	562	4.8	41
2	-20.62	181.03	650	4.2	15
3	-26.00	184.10	42	5.4	43
4	-17.97	181.66	626	4.1	19
5	-20.42	181.96	649	4.0	11
6	-19.68	184.31	195	4.0	12




- The boxplot of the earthquake magnitudes shows the output of summary except the mean, but it also shows *outliers*, extreme values that distort the mean:



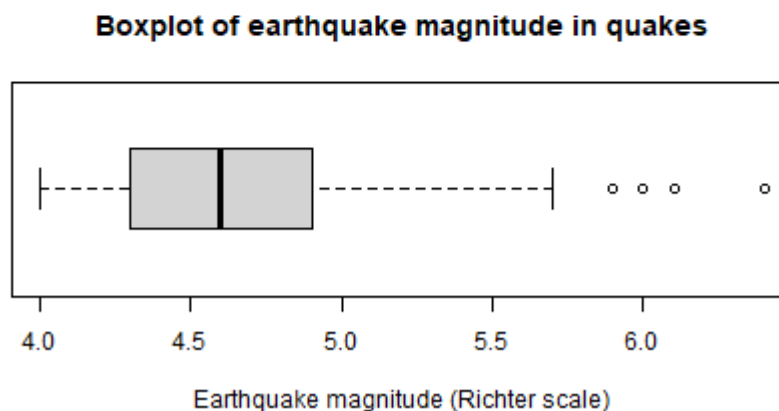
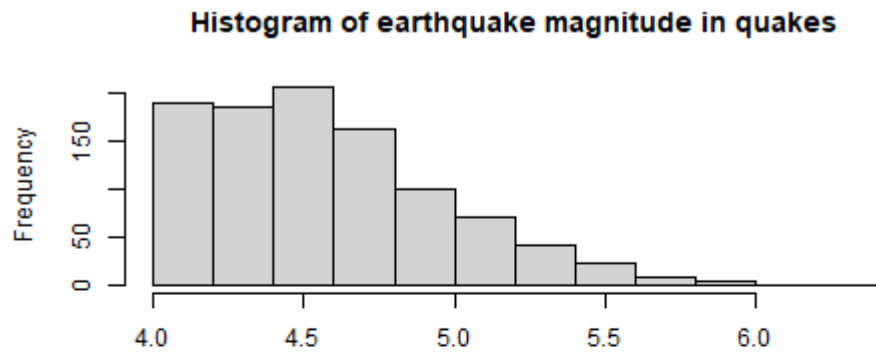
- Creating a boxplot is simple: for the earthquake magnitudes (`quakes$mag`):

```
boxplot(quakes$mag)
```



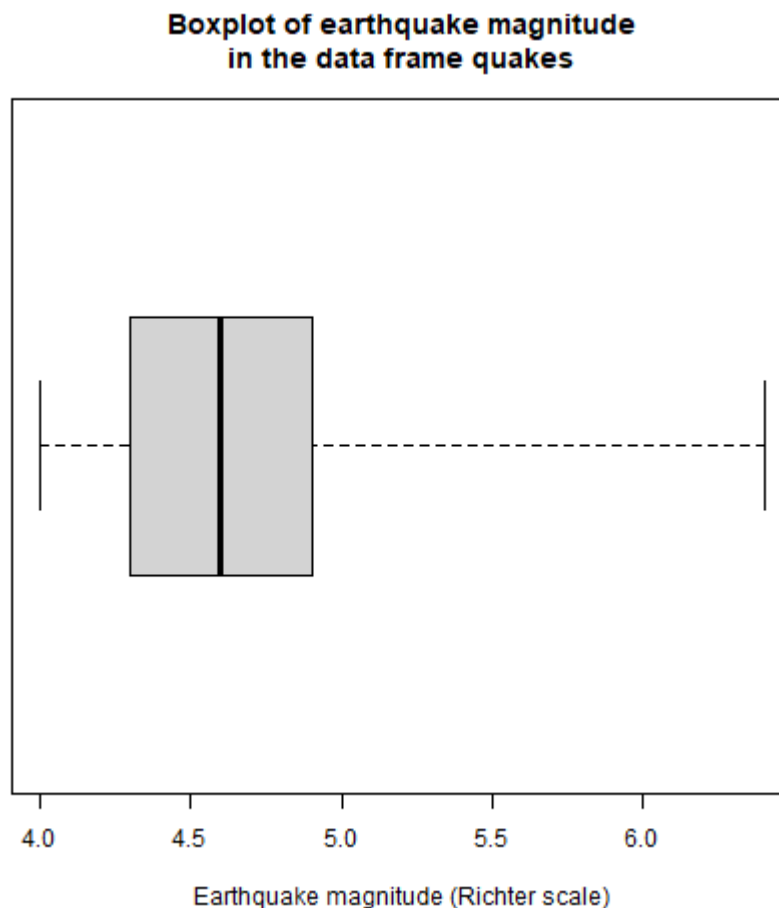
- Let's look at histogram and boxplot on top of one another, with a minimal customization:

```
par(mfrow=c(2,1))
hist(quakes$mag, xlab="",
     main="Histogram of earthquake magnitude in quakes")
boxplot(quakes$mag,
        horizontal=TRUE,
        main="Boxplot of earthquake magnitude in quakes",
        xlab="Earthquake magnitude (Richter scale)")
```



- Like the histogram, a boxplot shows important **global** (overall) features of the value distribution:
  - centrality (where is the midpoint of the distribution?)
  - spread (how far are the whiskers apart, how wide is the box?)
  - skewness (where is the box relative to the whiskers?)
- The boxplot does not show important **local** features, like
  - modes (multiple significant peaks or maxima)
  - valleys (local minima)
- Outliers are displayed explicitly (computed as 1.5 times the Inter-Quartile Range or IQR)
- The range parameter in boxplot determines how far the whiskers should extend from the box. range=0 includes all values.

```
boxplot(quakes$mag, horizontal=TRUE,
        range=0,
        main="Boxplot of earthquake magnitude\nin the data frame quakes",
        xlab="Earthquake magnitude (Richter scale)")
```



## 6 Side-by-side boxplots

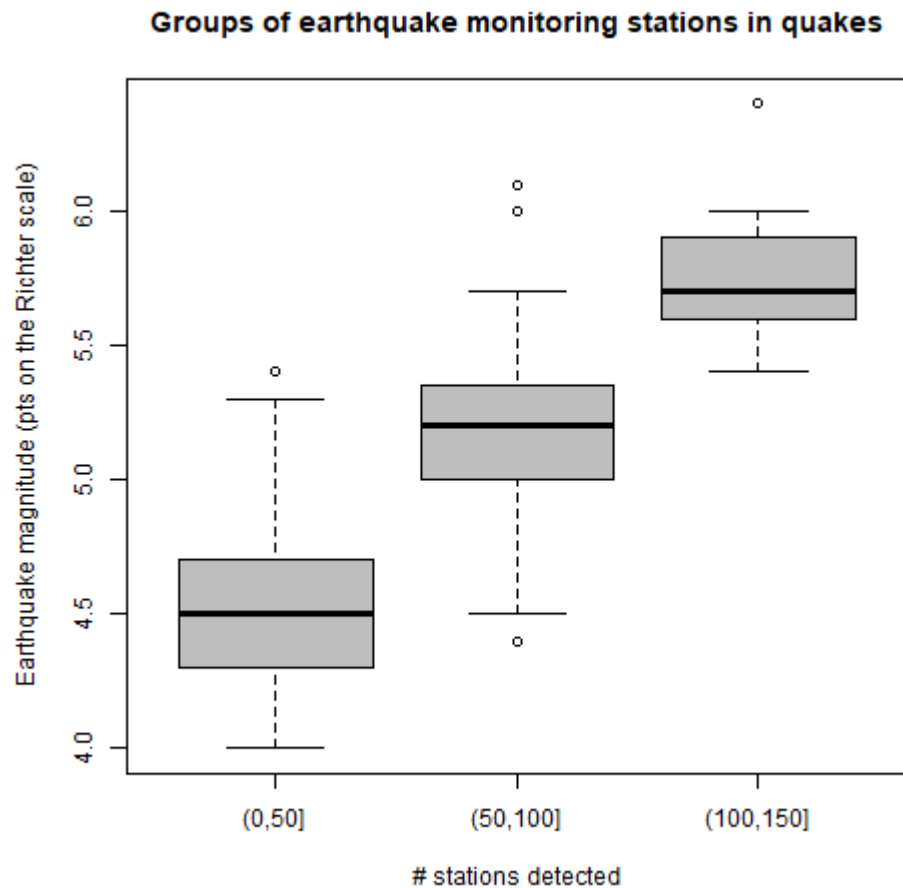
- Boxplots are useful to compare data features by plotting them side-by-side, e.g. for identifying how many monitoring stations detected each event
- We use `cut` to create three levels of stations for three boxes:

```
stations <- cut(quakes$stations, breaks=c(0,50,100,150))
stations[1:5] # first five elements
```

```
[1] (0,50] (0,50] (0,50] (0,50] (0,50]
Levels: (0,50] (50,100] (100,150]
```

- The factor `stations` breaks the observations in three groups
  - events detected by 50 stations or fewer `(0,50]`
  - events detected by 51 to 100 stations `(50,100]`
  - events detected by between 100 and 150 stations `(100,150]`
- The boxplot compares the distributions of the magnitudes of the events according to these three groups:

```
boxplot(quakes$mag ~ stations,
        main="Groups of earthquake monitoring stations in quakes",
        xlab="# stations detected",
        ylab="Earthquake magnitude (pts on the Richter scale)",
        col="gray")
```




- You can see that the higher the recorded magnitude, the more stations detected the given seismic event
- In the argument, we've used the *formula*  $y \sim x$  ("y vs. x") with  $x=\text{stations}$  and  $y=\text{quakes\$mag}$ . If the  $y \sim$  is missing, the argument is taken as  $x$  (cp. `help(boxplot)`).

## 7 Scatterplots


- Scatterplots are used to identify relationships between the observed values of two different numeric-continuous variables  $x, y$
- The scatterplot is displayed as an  $x$ - $y$ -coordinate plot but not every  $x$ - $y$ -plot shows relationships of interest - e.g. a plot of the latitude vs. longitude in quakes:

```
plot(x=quakes$lat,
     y=quakes$lon)
```




- A real scatterplot is the visualization of the mileage (mpg) vs. weight (wt) of cars in the built-in mtcars data set:

```
plot(mtcars$mpg ~ mtcars$wt)
title("Car mileage (mpg) vs. weight (wt) in mtcars")
```



## 8 Scatterplots of more than two variables

- The famous `iris` dataset, collected in the mid-1930s, contains petal and sepal measurements for three species of perennial iris flowers, *Iris setosa*, *Iris virginica*, and *Iris versicolor* (Fisher, 1936).



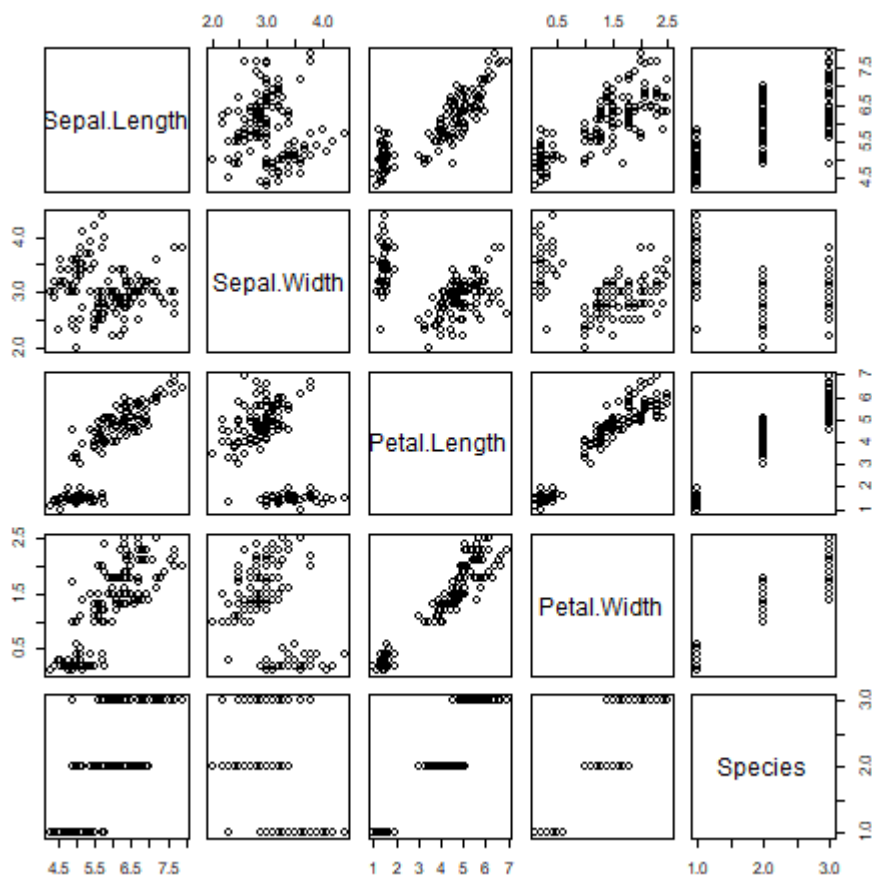
- View the first few records:

```
head(iris[1:5])
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa

- To plot every variable against every other, you can plot an array of x-y-plots:
  - Each column has the shown variable as x- and the others as y-axis
  - Each row has the shown variable as y- and the others as x-axis
  - E.g. the square (2,1) shows x = Sepal.Length, y = all others

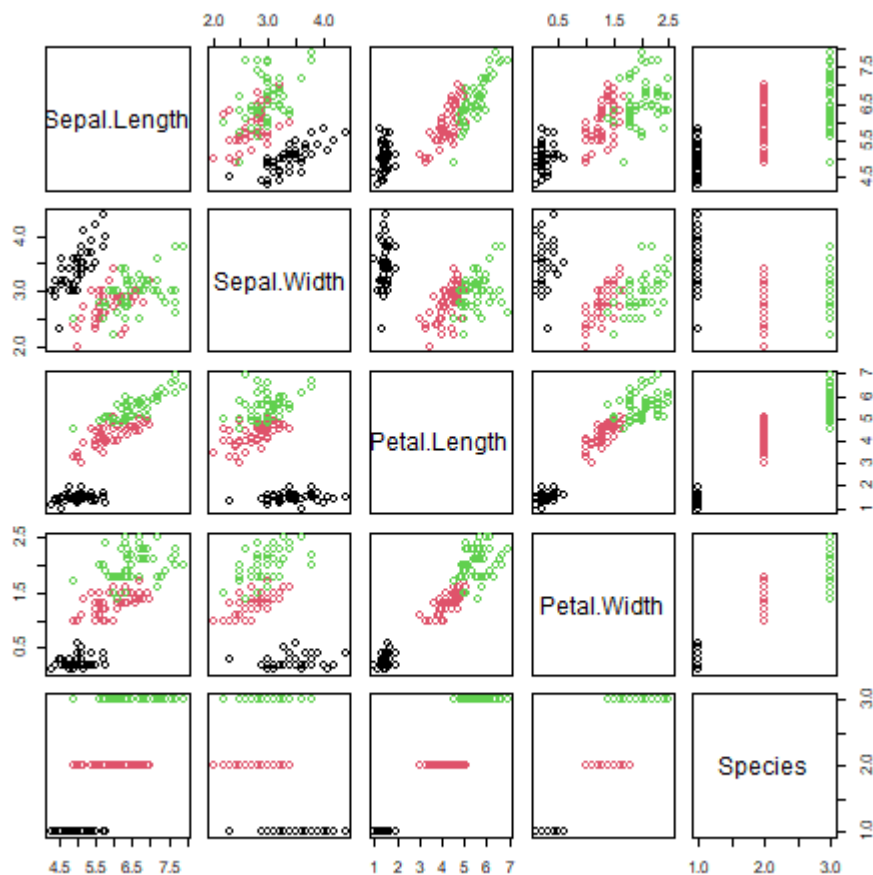
```
plot(iris)
```



- The array is much easier to read when adding the iris species as a sixth variable to the plot.




```
plot(iris, col=iris$Species)
```




- Homing in on one of the diagrams, e.g. Sepal.Width vs. Sepal.Length:

```
plot(
  iris$Sepal.Width ~ iris$Sepal.Length,
  col=iris$Species,
  pch=19)
```



- As an example for a fully customized plot with legend:

```
plot(
  y=iris$Petal.Length, xlab="Petal length (cm)",
  x=iris$Petal.Width, ylab="Petal width (cm)",
  col=iris$Species,
  pch=19)
legend("topleft",
  legend=c("setosa", "virginica", "versicolor"),
  col=c("black", "red", "green"), pch=19)
title("Petal width vs. length for three species of iris")
```



## 9 References

- ["ggplot2 barplots: Quick start guide", sthda.com](https://www.sthda.com/blog/ggplot2-barplots-quick-start-guide/)

Author: Marcus Birkenkrahe

Created: 2022-11-19 Sat 16:56